



Empolis Machine Learning Workshop

– März 2016 –

Features für Text

Prof. Dr. Adrian Ulges
Fachbereich DCSM / Informatik
Hochschule RheinMain

21. März 2016

Some Remarks regarding Text Features



ML Applications involving Text

- ▶ Information extraction / part-of-speech tagging
- ▶ Sentiment analysis
- ▶ Spam filtering
- ▶ Information retrieval
- ▶ Recommendation (of news, videos, movies, jobs, ...)

Some Remarks regarding Text Features



ML Applications involving Text

- ▶ Information extraction / part-of-speech tagging
- ▶ Sentiment analysis
- ▶ Spam filtering
- ▶ Information retrieval
- ▶ Recommendation (of news, videos, movies, jobs, ...)

Example: Bag-of-words Features

- ▶ Features for text should satisfy our criteria: **compactness**, **discriminativity**, **invariance**
- ▶ Common baseline (**bag-of-words** features): Store the count of appearances for each term in a document
- ▶ How would you **rank** bag-of-words features regarding compactness / discriminativity / invariance ?

Some Remarks regarding Text Features



Remarks

- ▶ In **this chapter**, we will have a look at some improvements over bag-of-words features
- ▶ The focus will still be on **simple text statistics**
- ▶ A very useful reference: Python's `nltk` module!

Text Features: Segmentation



- ▶ First Question: What is a “**term**”?
- ▶ **Text segmentation** into terms is not a trivial problem
(*splitting by spaces does not entirely do the job!*)

Text Features: Segmentation



- ▶ First Question: What is a “**term**”?
- ▶ **Text segmentation** into terms is not a trivial problem
(*splitting by spaces does not entirely do the job!*)

Example	Approach
Germany's chancellor	<i>rule-based recognition</i>
3/20/91 vs. Mar 12, 1991	<i>rule-based recognition</i>
(0049) 611/9495-1215	<i>rule-based recognition</i>
San Francisco	

Text Features: Segmentation



- ▶ First Question: What is a “**term**”?
- ▶ **Text segmentation** into terms is not a trivial problem
(*splitting by spaces does not entirely do the job!*)

Example	Approach
Germany's chancellor	<i>rule-based recognition</i>
3/20/91 vs. Mar 12, 1991	<i>rule-based recognition</i>
(0049) 611/9495-1215	<i>rule-based recognition</i>
San Francisco	<i>statistical methods</i>
Lebensversicherungsgesellschaft vs. Malerei	

Text Features: Segmentation



- ▶ First Question: What is a “**term**”?
- ▶ **Text segmentation** into terms is not a trivial problem
(*splitting by spaces does not entirely do the job!*)

Example	Approach
Germany's chancellor	<i>rule-based recognition</i>
3/20/91 vs. Mar 12, 1991	<i>rule-based recognition</i>
(0049) 611/9495-1215	<i>rule-based recognition</i>
San Francisco	<i>statistical methods</i>
Lebensversicherungsgesellschaft vs. Malerei	<i>compound splitter (dictionary- based vs. statistical methods)</i>

Text Features: Segmentation

Code Example: Python

- ▶ This code uses **regular expressions**, which allow us to search a wide range of text patterns in strings

Operator	Behavior
.	Wildcard, matches any character
^abc	Matches some pattern <i>abc</i> at the start of a string
abc\$	Matches some pattern <i>abc</i> at the end of a string
[abc]	Matches one of a set of characters
[A-Z0-9]	Matches one of a range of characters
ed ing s	Matches one of the specified strings (disjunction)
*	Zero or more of previous item, e.g., <i>a*</i> , <i>[a-z]*</i> (also known as <i>Kleene Closure</i>)
+	One or more of previous item, e.g., <i>a+</i> , <i>[a-z]+</i>
?	Zero or one of the previous item (i.e., optional), e.g., <i>a?</i> , <i>[a-z]?</i>
{n}	Exactly <i>n</i> repeats where <i>n</i> is a non-negative integer
{n,}	At least <i>n</i> repeats
{,n}	No more than <i>n</i> repeats
{m,n}	At least <i>m</i> and no more than <i>n</i> repeats
a(b c)+	Parentheses that indicate the scope of the operators

```
>>> text = 'That U.S.A. poster-print costs $12.40...'
>>> pattern = r'^(?x)      # set flag to allow verbose regexps
...      ([A-Z]\.)+        # abbreviations, e.g. U.S.A.
...      | \w+(-\w+)*       # words with optional internal hyphens
...      | \$?\d+(\.\d+)?%?  # currency and percentages, e.g. $12.40, 82%
...      | \.\.\.           # ellipsis
...      | [[\.,;"'()?():-_\]] # these are separate tokens
...
>>> nltk.regexp_tokenize(text, pattern)
['That', 'U.S.A.', 'poster-print', 'costs', '$12.40', '...']
```

Text Features: Normalization



- ▶ We also **normalize** text to increase robustness to flexion and sentence structure
- ▶ **Step 1:** Lower-casing (*Sometimes* → *sometimes*)
- ▶ **Step 2:** Stemming = reducing words to their stem

Text Features: Normalization



- ▶ We also **normalize** text to increase robustness to flexion and sentence structure
- ▶ **Step 1:** Lower-casing (*Sometimes* → *sometimes*)
- ▶ **Step 2:** Stemming = reducing words to their stem

Stemming: Methods

- ▶ Rule-based Methods
 - ▶ Example rule: $*t \rightarrow *$ (geht → geh)
 - ▶ Example rule: $*en \rightarrow *$ (gehen → geh)
- ▶ Dictionary-based Methods
 - ▶ Example: `stem['ging'] = 'geh'`
 - ▶ popular for languages with strong flexion (*like German*)

Stemming: Code Example



```
1  def naive_stem(word):
2      regexp = r'^(.*?)(ing|ly|ed|ious|ies|ive|es|s|ment)?'
3      stem, suffix = re.findall(regexp, word)[0]
4      return stem
5
6  >>> tokens = ['women', 'swords', 'is', 'lying']
7
8  >>> [naive_stem(t) for t in tokens]
9
10     ['women', 'sword', 'i', 'ly']           // naive
```

Stemming: Code Example



```
1  def naive_stem(word):
2      regexp = r'^(.*?)(ing|ly|ed|ious|ies|ive|es|s|ment)?'
3      stem, suffix = re.findall(regexp, word)[0]
4      return stem
5
6  >>> tokens = ['women', 'swords', 'is', 'lying']
7
8  >>> [naive_stem(t) for t in tokens]
9
10     ['women', 'sword', 'i', 'ly']           // naive
11
12  >>> [nltk.WordNetLemmatizer().lemmatize(t)
13       for t in tokens]
14
15     ['woman', 'sword', 'is', 'lying']       // dict-based
16
17  >>> [nltk.PorterStemmer().stem(t)
18       for t in tokens]
19
20     ['women', 'sword', 'is', 'lie']         // rule-based
21
```

Text Features: Synsets

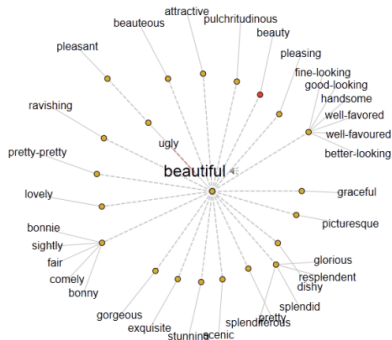


- Can we achieve invariance to **synonyms**?

"What a beautiful day!" vs.

"What a lovely day!"

- A frequent approach are **thesauri**: A thesaurus is a collection of terms, connected by (pre-defined) relations



Text Features: Synsets

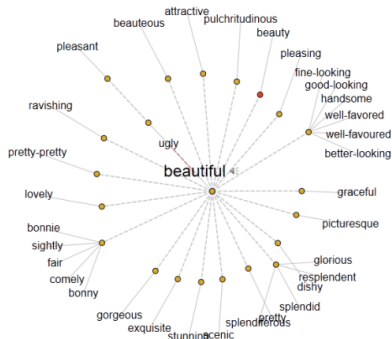


- ▶ Can we achieve invariance to **synonyms**?

"What a beautiful day!" vs.

"What a lovely day!"

- ▶ A frequent approach are **thesauri**: A thesaurus is a collection of terms, connected by (pre-defined) relations
- ▶ Typical relations
 - ▶ synonyms (*beautiful vs. lovely*)
 - ▶ antonym (*beautiful vs. ugly*)
 - ▶ generalization/specialization (*a boat is a vehicle*)
- ▶ Synonyms form so-called **synsets**



Synsets: Python Example



```
1  >>> from nltk.corpus import wordnet as wn
2  >>> wn.synsets("dog")
3
4  [Synset('dog.n.01'),
5   Synset('frump.n.01'),
6   Synset('dog.n.03'),
7   Synset('cad.n.01'),
8   Synset('frank.n.02'),
9   Synset('pawl.n.01'),
10  Synset('andiron.n.01'),
11  Synset('chase.v.01')]
```


Synsets: Python Example



```
>>> from nltk.corpus import wordnet as wn
>>> wn.synsets("dog")

[Synset('dog.n.01'),
 Synset('frump.n.01'),
 Synset('dog.n.03'),
 Synset('cad.n.01'),
 Synset('frank.n.02'),
 Synset('pawl.n.01'),
 Synset('andiron.n.01'),
 Synset('chase.v.01')]

>>> for synset in wn.synsets("dog"):
    print "dog =", synset.definition

dog = a member of the genus Canis ...
dog = a dull unattractive unpleasant woman
dog = informal term for a man
dog = a smooth-textured sausage ...
dog = metal supports for logs in a fireplace
dog = go after with the intent to catch
...
```

From Thesauri to Ontologies



- ▶ We can extend the concept of a thesaurus to *ontologies*
- ▶ An ontology can be thought of as a generalized **knowledge base** containing objects and relations between them
- ▶ Ontologies can be **combined** by linking their objects

```
{{Infobox Town AT |
name = Innsbruck |
image_coa = InnsbruckWappen.png |
image_map = Karte-tirol-I.png |
state = {{Tyrol}} |
regbszk = {{Statutory city}} |
population = 117,342 |
population_as_of = 2006 |
pop_dens = 1,119 |
area = 104.91 |
elevation = 574 |
lat_deg = 47 |
lat_min = 16 |
lat_sec = N |
lon_deg = 11 |
lon_min = 23 |
lon_sec = E |
postal_code = 6010-6060 |
area_code = 0512 |
licence = I |
mayor = Hilde Zach |
website = (http://innsbruck.at) |
}}
```

Innsbruck	
	
Country	 Austria
State	Tyrol
Administrative region	Statutory city
Population	117,342 (2006)
Area	104.91 km²
Population density	1,119 /km²
Elevation	574 m
Coordinates	 47°16′ N 11°23′ E
Postal code	6010-6060
Area code	0512
Licence plate code	I
Mayor	Hilde Zach
Website	www.innsbruck.at

From Thesauri to Ontologies



- ▶ We can extend the concept of a thesaurus to *ontologies*
- ▶ An ontology can be thought of as a generalized **knowledge base** containing objects and relations between them
- ▶ Ontologies can be **combined** by linking their objects

Example: The DBPedia Project

- ▶ 20.8 mio. "things", crawled from Wikipedia infoboxes
- ▶ > 500 mio. "facts"
- ▶ representation by RDF (*Resource Description Framework*)

```
{(Infobox Town AT |  
  name = Innsbruck |  
  image_coa = InnsbruckWappen.png |  
  image_map = Karte-tirol-I.png |  
  state = {{Tyrol}} |  
  regbrk = {{Statutory city}} |  
  population = 117,342 |  
  population_as_of = 2008 |  
  pop_dens = 1,119 |  
  area = 104.91 |  
  elevation = 574 |  
  lat_deg = 47 |  
  lat_min = 16 |  
  lat_sec = N |  
  lon_deg = 11 |  
  lon_min = 23 |  
  lon_sec = E |  
  postal_code = 6010-6080 |  
  area_code = 0512 |  
  licence = I |  
  mayor = Hilde Zach |  
  website = (http://innsbruck.at) |  
}}
```

Innsbruck	
	
Country	Austria
State	Tyrol
Administrative region	Statutory city
Population	117,342 (2008)
Area	104.91 km²
Population density	1,119 /km²
Elevation	574 m
Coordinates	47°16' N 11°23' E @
Postal code	6010-6080
Area code	0512
Licence plate code	I
Mayor	Hilde Zach
Website	www.innsbruck.at @

From Thesauri to Ontologies



- ▶ We can extend the concept of a thesaurus to *ontologies*
- ▶ An ontology can be thought of as a generalized **knowledge base** containing objects and relations between them
- ▶ Ontologies can be **combined** by linking their objects

Example: The DBPedia Project

- ▶ 20.8 mio. "things", crawled from Wikipedia infoboxes
- ▶ > 500 mio. "facts"
- ▶ representation by RDF (*Resource Description Framework*)
- ▶ allows **smarter search** ("give me all cities in New Jersey with more than 10,000 inhabitants")

```
{{Infobox Town AT |
name      = Innsbruck |
image_coa = InnsbruckWappen.png |
image_map = Karte-tirol-1.png |
state     = {{Tyrol}} |
regbrk    = {{{Statutory city}}} |
population      = 117,342 |
population_as_of      = 2008 |
pop_dens      = 1,119 |
area      = 104.91 |
elevation      = 574 |
lat_deg      = 47 |
lat_min      = 16 |
lat_sec      = N |
lon_deg      = 11 |
lon_min      = 23 |
lon_sec      = E |
postal_code      = 6010-6080 |
area_code      = 0512 |
licence      = I |
mayor      = Hilde Zach |
website      = {http://innsbruck.at}
}}
```

Innsbruck	
	
Country	 Austria
State	Tyrol
Administrative region	Statutory city
Population	117,342 (2008)
Area	104.91 km²
Population density	1,119 /km²
Elevation	574 m
Coordinates	 47°16′ N 11°23′ E
Postal code	6010-6080
Area code	0512
Licence plate code	I
Mayor	Hilde Zach
Website	www.innsbruck.at

Text Features: N-Grams



- So far, we have neglected the **order** of words in the document

*"I can **not** believe it – **What** a **cool** video!" vs.*

*"This video is **not cool** – **What** a..."*

Text Features: N-Grams



- ▶ So far, we have neglected the **order** of words in the document
 *"I can **not** believe it – **What** a **cool** video!" vs.*
 *"This video is **not cool** – **What** a..."*
- ▶ A simple statistical approach are **n-grams**: Instead of segmenting text into single tokens, we segment it into subsequences of n tokens each!

Text Features: N-Grams



- ▶ So far, we have neglected the **order** of words in the document
 *"I can **not** believe it – **What** a **cool** video!" vs.*
 *"This video is **not cool** – **What** a..."*
- ▶ A simple statistical approach are **n-grams**: Instead of segmenting text into single tokens, we segment it into subsequences of n tokens each!

In the Example

- ▶ bag-of-words feature

$$\left\{ (This: 1), (video: 1), (is: 1), (cool: 1), \dots \right\}$$

- ▶ n-gram feature

$$\left\{ (This\ video: 1), (video\ is: 1), (is\ not: 1), (not\ cool: 1), \dots \right\}$$

Text Features: N-Grams



- ▶ So far, we have neglected the **order** of words in the document
*"I can **not** believe it – **What** a **cool** video!"* vs.
*"This video is **not cool** – **What** a..."*
- ▶ A simple statistical approach are **n-grams**: Instead of segmenting text into single tokens, we segment it into subsequences of n tokens each!

In the Example

- ▶ bag-of-words feature

$$\left\{ (This: 1), (video: 1), (is: 1), (cool: 1), \dots \right\}$$

- ▶ n-gram feature

$$\left\{ (This\ video: 1), (video\ is: 1), (is\ not: 1), (not\ cool: 1), \dots \right\}$$

- ▶ Problem: Features get (even more) **high-dimensional**!